

EPBI 414

Unit 7

Advanced SQL

The Recap - Unit 6

- SQL and MySQL
 - Connecting using the command line interface, MySQL Workbench
- Coding standards and conventions
- Basics of SQL
 - Selecting, filtering, inner joins
 - Symbolic logic

Unit 7 Overview

- More complex joins
 - RIGHT / LEFT JOIN, FULL JOIN
 - ON **vs.** WHERE
- UNION, UNION ALL
- Data Summary and Analysis
- Modifying Data

More complicated joins

- Last time, we covered the *inner join*
 - Returns rows when the criteria is met
- The *left / right outer join* returns all records in one table, and matches from the other
 - The *outer* is usually omitted
- The *full join* (sometimes *full outer join*) gives you all records from both tables
 - Not supported in MySQL

An example LEFT JOIN

```
SELECT p.patient_id,  
       p.dob,  
       v.visit_id,  
       v.provider_id,  
       v.visit_date  
FROM patients AS p  
LEFT JOIN last_visit AS v  
ON p.patient_id = v.patient_id;
```

patient_id	pcp_id	dob
1	1	12/08/1944
2	1	08/03/1999
3	2	02/07/1983

visit_id	patient_id	provider_id	visit_date
1	1	3	02/12/2016
2	3	1	09/03/2016

LEFT JOIN Results

```
SELECT p.patient_id,  
       p.dob,  
       v.visit_id,  
       v.provider_id,  
       v.visit_date  
FROM patients AS p  
LEFT JOIN last_visit AS v  
ON p.patient_id = v.patient_id;
```

patient_id	dob	visit_id	provider_id	visit_date
1	12/08/1944	1	3	02/12/2016
2	08/03/1999	NULL	NULL	NULL
3	02/07/1983	2	1	09/03/2016

ON VS. WHERE

- Left & right joins give us an opportunity to discuss the difference between `ON` and `WHERE`
- In short: `WHERE` filters results, `ON` joins rows together
- What's the difference?

WHERE filters rows

```
SELECT p.patient_id,  
       p.dob,  
       v.visit_id,  
       v.provider_id,  
       v.visit_date  
FROM patients AS p  
LEFT JOIN last_visit AS v  
ON p.patient_id = v.patient_id  
WHERE p.dob >= '1950-01-01';
```

patient_id	dob	visit_id	provider_id	visit_date
2	08/03/1999	NULL	NULL	NULL
3	02/07/1983	2	1	09/03/2016

ON joins rows together

```
SELECT p.patient_id,  
       p.dob,  
       v.visit_id,  
       v.provider_id,  
       v.visit_date  
FROM patients AS p  
LEFT JOIN last_visit AS v  
ON p.patient_id = v.patient_id  
   AND p.dob >= '1950-01-01';
```

patient_id	dob	visit_id	provider_id	visit_date
1	12/08/1944	NULL	NULL	NULL
2	08/03/1999	NULL	NULL	NULL
3	02/07/1983	2	1	09/03/2016

ON and inner vs. outer

- Inner joins only return a row when the ON condition is met
 - So, ON and WHERE are equivalent!
- Outer joins always return some rows
 - This makes ON and WHERE different!
- What is the practical meaning of this?

Writing good queries

- It's tempting to think something like...
 - "If the `ON` controls which rows get included in the join at all, while the `WHERE` filters out rows after the joining is done...I should use `ON` to return smaller sets of data to start with!"
- Unfortunately, this instinct doesn't quite work
- The "sequence" isn't that simple, because...

The query optimizer

- Every RDBMS has something called a *query optimizer*
- Query optimization varies between database engines
 - Fairly low-level technical part of the RDBMS
- The optimizer makes "hard and fast" rules difficult to articulate

EXPLAIN

- One neat feature of most RDBMSs is the `EXPLAIN` feature
- By running `EXPLAIN <QUERY>`, you can get your RDBMS to tell you its plan
- This can let you (more likely, expert DBAs) identify bottlenecks and improve queries

```
λ Cmder
mysql> EXPLAIN
      -> SELECT * FROM departments;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | type | possible_keys | key      | key_len | ref | rows | Extra      |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE      | departments | index | NULL          | dept_name | 42      | NULL | 9 | Using index |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql> |
```


Example EXPLAIN Query

EXPLAIN and You

- EXPLAIN is an in-depth RDBMS feature
 - Also, very dependant on your RDBMS
- The query optimizer isn't magic
 - Can't save a terrible query from itself
- You should focus on logical queries that are efficient ***enough***
 - "Premature optimization is the root of all evil."
Attributed to Donald Knuth

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
HOW MUCH TIME YOU SHAVE OFF	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS



Two relevant xkcd comics^{1,2}

A rule of thumb

- Keeping in mind that "hard and fast" rules are hard to come by...
- When you are writing your query, you can think of it like this:

ON ***happens before*** WHERE
or
Joins before filters

FULL JOIN

- As mentioned, MySQL does not support ***full outer joins*** (or just ***full joins***)
- Full joins are relatively uncommon
 - There are few circumstances where you want all the rows from both tables, regardless of match
- You can simulate a full join in MySQL if you need it
 - Not required for this class

Other join notes

- You can join a table against itself - this is called a ***self join***
 - Sometimes, it can make filtering faster - really depends on the optimizer
- Many RDBMSs support other join types
 - Natural join, cross join, implicit join...
- For your day-to-day, stick to the main ones

Many-to-many joins

- Caution is advised when joining on two columns which both contain duplicates
- SQL will produce all possible combinations from your join query
 - This may not be - probably isn't - what you want
- Many-to-many is also expensive computationally

Example many-to-many

```
SELECT  s.*,
        t.*
FROM
(
  SELECT  *
          FROM salaries
          WHERE emp_no = 10004
) AS s
INNER JOIN
(
  SELECT  *
          FROM titles
          WHERE emp_no = 10004
) AS t
ON s.emp_no = t.emp_no;
```

UNION

- The `UNION` command allows you to "stack" the results of two queries vertically
- Useful when you run two generally similar queries with different filters
- Columns need to line up, and be compatible!
 - The RDBMS can usually **coerce** values from one type to another (i.e. numeric to character)

UNION Syntax

- Unions are performed by writing `SELECT` statements, separated by `UNION`

```
SELECT * FROM tbl1  
UNION  
SELECT * FROM tbl2;
```

UNION vs. UNION ALL

- The `UNION` statement *often* (no guarantees!) removes duplicate rows
- If this is not desired behavior, `UNION ALL` will generally override it
- This tends to depend on your RDBMS

Summarizing Data

- So far, we have just accessed and filtered data using SQL
- But SQL can do much more
 - Aggregation
 - Grouping
 - Calculations
- These features are often very fast, compared to other software

SQL Analysis Tools

- Calculated Fields
 - SQL functions
- Selecting unique records (`DISTINCT`)
- Grouping records together (`GROUP BY`)
 - Aggregation functions

Getting calculated fields

- Most of the time, a ***field*** is synonymous with a ***column***
 - Codd never talked about ***fields*** at all
- However, most RDBMSs will let you add ***calculated fields*** to your queries
- Client apps don't see any difference between ***calculated fields*** and regular columns

Constant value example

```
SELECT 'A' AS a;
```

```
mysql> SELECT 'A' AS a;
+----+
| a  |
+----+
| A  |
+----+
1 row in set (0.00 sec)

mysql> |
```

Naming calculations

- You can name your calculated fields using the `AS` statement
 - Just like subqueries
- Calculated field names tend not to be available within the same query
 - Can be accessed from subqueries
- Very useful if you are using `UNION`!

Using SQL Functions

- SQL allows you to use both basic mathematics and built-in SQL functions in your calculated fields
- There are a *lot* of SQL functions
 - Can't possibly cover them all in class
- Here is an example: **concatenating** two strings together

Example SQL function

```
SELECT CONCAT(`first_name`,` `,`last_name`) AS full_name,  
        hire_date  
FROM employees  
WHERE last_name LIKE 'Terw%'  
LIMIT 3;
```

```
mysql> SELECT CONCAT(`first_name`,` `,`last_name`) AS full_name,  
->        hire_date  
->        FROM employees  
->        WHERE last_name LIKE 'Terw%'  
->        LIMIT 3;
```

full_name	hire_date
Suebskul Terwilliger	1995-12-29
Shounak Terwilliger	1989-05-05
Marsha Terwilliger	1988-03-04

3 rows in set (0.00 sec)

```
mysql>
```

Another function

```
SELECT emp_no,  
       first_name,  
       last_name,  
       TIMESTAMPDIFF(year,birth_date,hire_date) AS age_at_hire  
FROM employees LIMIT 10;
```

```
mysql> SELECT emp_no,  
->         first_name,  
->         last_name,  
->         TIMESTAMPDIFF(year,birth_date,hire_date) AS age_at_hire  
->         FROM employees LIMIT 10;
```

emp_no	first_name	last_name	age_at_hire
10001	Georgi	Facello	32
10002	Bezalel	Simmel	21
10003	Parto	Bamford	26
10004	Chirstian	Koblick	32
10005	Kyoichi	Maliniak	34
10006	Anneke	Preusig	36
10007	Tzvetan	Zielinski	31
10008	Saniya	Kalloufi	36
10009	Sumant	Peac	32
10010	Duangkaew	Piveteau	26

```
10 rows in set (0.00 sec)  
  
mysql> |
```


Using field names - bad

```
SELECT emp_no,  
       emp_no + 10 AS emp_no_plus,  
       emp_no_plus * emp_no_plus AS emp_no_plus_squared  
FROM employees  
LIMIT 10;
```

```
mysql> SELECT emp_no,  
->         emp_no + 10 AS emp_no_plus,  
->         emp_no_plus * emp_no_plus AS emp_no_plus_squared  
->         FROM employees  
->         LIMIT 10;  
ERROR 1054 (42S22): Unknown column 'emp_no_plus' in 'field list'  
mysql> |
```

Using field names - good

```
SELECT subq.emp_no,  
       subq.emp_no_plus,  
       subq.emp_no_plus * subq.emp_no_plus AS emp_no_plus_squared  
FROM  
(  
  SELECT emp_no,  
         emp_no + 10 AS emp_no_plus  
    FROM employees  
) AS subq  
LIMIT 10;
```

```
mysql> SELECT subq.emp_no,  
->         subq.emp_no_plus,  
->         subq.emp_no_plus * subq.emp_no_plus AS emp_no_plus_squared  
->       FROM  
->       (  
->         SELECT emp_no,  
->                emp_no + 10 AS emp_no_plus  
->           FROM employees  
->       ) AS subq  
->       LIMIT 10;
```

emp_no	emp_no_plus	emp_no_plus_squared
10001	10011	100220121
10002	10012	100240144
10003	10013	100260169
10004	10014	100280196
10005	10015	100300225
10006	10016	100320256
10007	10017	100340289
10008	10018	100360324
10009	10019	100380361
10010	10020	100400400

10 rows in set (0.11 sec)

```
mysql>
```

DISTINCT

- Often, you have tables with values repeated across rows
- `DISTINCT` is a keyword that helps us obtain a unique list of values
- Often, this is useful when counting (coming up)

DISTINCT values

```
SELECT emp_no  
FROM salaries  
LIMIT 10;
```

```
mysql> SELECT emp_no  
-> FROM salaries  
-> LIMIT 10;  
  
+-----+  
| emp_no |  
+-----+  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
| 10001 |  
+-----+  
10 rows in set (0.00 sec)  
  
mysql>
```

```
SELECT DISTINCT emp_no  
FROM salaries  
LIMIT 10;
```

```
mysql> SELECT DISTINCT emp_no  
-> FROM salaries  
-> LIMIT 10;  
  
+-----+  
| emp_no |  
+-----+  
| 10001 |  
| 10002 |  
| 10003 |  
| 10004 |  
| 10005 |  
| 10006 |  
| 10007 |  
| 10008 |  
| 10009 |  
| 10010 |  
+-----+  
10 rows in set (0.00 sec)  
  
mysql> |
```

DISTINCT on rows

- You can use the `DISTINCT` keyword on multiple variables
- MySQL will return the distinct tuples from your variables
- Cannot easily return `DISTINCT` plus some non-distinct value
 - That's where grouping comes in

Caveats on DISTINCT

- General caveat: by definition, a SQL table (relation) is an unordered set
 - If you don't use an `ORDER BY`, you have no guarantee that the order will be consistent between queries
- Using `DISTINCT` on multiple columns can be very time-consuming
 - Indexes can help

Break Time

Grouping and Aggregation

- One of the most powerful tools that you have in SQL is the ability to ***group*** and ***aggregate*** data
- By using the five basic SQL aggregation functions, and a `GROUP BY` statement, you can get very useful information

Aggregation Functions

- SQL has five basic aggregation functions:
 - `AVG ()`
 - `COUNT ()`
 - `MAX ()`
 - `MIN ()`
 - `SUM ()`

Notes on aggregations

- The aggregation functions only work either in *isolation*, or as part of a *grouping*
- You can't ask for a regular column and a aggregation without doing a group
- You might not be able to use certain aggregations on certain data types

Basic aggregation

```
SELECT AVG(salary) AS average_salary  
FROM salaries;
```

```
mysql> SELECT AVG(salary) AS average_salary  
->      FROM salaries;
```

```
+-----+  
| average_salary |  
+-----+  
|      63810.7448 |  
+-----+  
1 row in set (0.63 sec)
```

The average of all salary records
in company history...(not super useful)

Basic frequencies

```
SELECT  title,
        gender,
        COUNT(*) AS n_employees FROM
        (
        SELECT  e.emp_no,
                e.gender,
                t.title
        FROM employees AS e
        INNER JOIN
        (
        SELECT  emp_no,
                title
        FROM titles
        WHERE to_date = '9999-01-01'
        ) AS t
        ON e.emp_no = t.emp_no
        ) AS subq
GROUP BY title,gender;
```

```

mysql> SELECT title,
-> gender,
-> COUNT(*) AS n_employees FROM
-> (
-> SELECT e.emp_no,
-> e.gender,
-> t.title
-> FROM employees AS e
-> INNER JOIN
-> (
-> SELECT emp_no,
-> title
-> FROM titles
-> WHERE to_date = '9999-01-01'
-> ) AS t
-> ON e.emp_no = t.emp_no
-> ) AS subq
-> GROUP BY title,gender;

```

title	gender	n_employees
Assistant Engineer	M	2148
Assistant Engineer	F	1440
Engineer	M	18571
Engineer	F	12412
Manager	M	5
Manager	F	4
Senior Engineer	M	51533
Senior Engineer	F	34406
Senior Staff	M	49232
Senior Staff	F	32792
Staff	M	15436
Staff	F	10090
Technique Leader	M	7189
Technique Leader	F	4866

14 rows in set (0.65 sec)

mysql>

How the frequencies look

HAVING vs. WHERE

- When you aggregate, `WHERE` filters out rows *before* the calculations
- To filter on the resulting groups, you use `HAVING` instead
 - Or you can put your aggregation in a subquery...
- The performance implications are generally subtle - and relate back to `EXPLAIN`

Another rule of thumb

- Another general rule of thumb is:

Use WHERE **before** GROUP BY

Use HAVING **after** GROUP BY

A HAVING example

```
SELECT  title,
        gender,
        COUNT(*) AS n_employees FROM
        (
        SELECT  e.emp_no,
                e.gender,
                t.title
        FROM employees AS e
        INNER JOIN
        (
        SELECT  emp_no,
                title
        FROM titles
        WHERE to_date = '9999-01-01'
        ) AS t
        ON e.emp_no = t.emp_no
        ) AS subq
GROUP BY title,gender
HAVING gender = 'M';
```



```

mysql> SELECT title,
-> gender,
-> COUNT(*) AS n_employees FROM
-> (
-> SELECT e.emp_no,
-> e.gender,
-> t.title
-> FROM employees AS e
-> INNER JOIN
-> (
-> SELECT emp_no,
-> title
-> FROM titles
-> WHERE to_date = '9999-01-01'
-> ) AS t
-> ON e.emp_no = t.emp_no
-> ) AS subq
-> GROUP BY title,gender
-> HAVING gender = 'M';

```

title	gender	n_employees
Assistant Engineer	M	2148
Engineer	M	18571
Manager	M	5
Senior Engineer	M	51533
Senior Staff	M	49232
Staff	M	15436
Technique Leader	M	7189

7 rows in set (0.67 sec)

HAVING Version

A WHERE example

```
SELECT  title,
        gender,
        COUNT(*) AS n_employees FROM
    (
        SELECT  e.emp_no,
                e.gender,
                t.title
        FROM employees AS e
        INNER JOIN
        (
            SELECT  emp_no,
                    title
            FROM titles
            WHERE to_date = '9999-01-01'
        ) AS t
        ON e.emp_no = t.emp_no
    ) AS subq
WHERE gender = 'M'
GROUP BY title,gender;
```

```

mysql> SELECT title,
-> gender,
-> COUNT(*) AS n_employees FROM
-> (
-> SELECT e.emp_no,
-> e.gender,
-> t.title
-> FROM employees AS e
-> INNER JOIN
-> (
-> SELECT emp_no,
-> title
-> FROM titles
-> WHERE to_date = '9999-01-01'
-> ) AS t
-> ON e.emp_no = t.emp_no
-> ) AS subq
-> WHERE gender = 'M'
-> GROUP BY title,gender;

```

title	gender	n_employees
Assistant Engineer	M	2148
Engineer	M	18571
Manager	M	5
Senior Engineer	M	51533
Senior Staff	M	49232
Staff	M	15436
Technique Leader	M	7189

7 rows in set (0.61 sec)

WHERE Version

DISTINCT COUNTing

```
SELECT COUNT(DISTINCT emp_no) AS n_employees  
FROM salaries  
WHERE to_date = '9999-01-01';
```

```
mysql> SELECT COUNT(DISTINCT emp_no) AS n_employees  
->      FROM salaries  
->      WHERE to_date = '9999-01-01';  
+-----+  
| n_employees |  
+-----+  
|      240124 |  
+-----+  
1 row in set (0.83 sec)  
  
mysql>
```

Putting it together

- SQL has many individual tools that are useful
 - Grouping, aggregating, ordering, limiting, filtering, et cetera
- You can combine these tools to answer relatively specific questions!
- You will learn to "think SQL"

Modifying data

- So far, all we've discussed has been pulling data out of SQL
- To close, we'll touch upon some basics of creating and modifying your own tables
- Unlike before, you will have destructive powers now
 - Use them wisely

Note

- Though we're going to cover some basics of database administration and modifying data, the course focus is on ***consumption***
- This is a very broad survey of these topics
- Does not substitute for having experience or a good DBA

What types of powers?

- Adding, modifying, and deleting data from existing tables
- Creating, modifying, and deleting tables
 - Setting data types, learning about constraints
 - We won't get into creating and deleting schema
- Using temporary tables, if you can
 - A very useful feature!

INSERT UPDATE DELETE

- INSERT is used to add new records to a table
- UPDATE is used to change the values of some rows of a table
- DELETE is used to remove rows from a table

INSERTing yourself

- Data insertion is generally the least risky of operations
 - Not risk free
- Generally, the syntax is:

```
INSERT INTO sometable (somecol1,somecol2,somecol3)  
VALUES (value1,value2,value3);
```

```
INSERT INTO sometable VALUES (value1,value2,value3);
```

Making UPDATES

- UPDATE is a more dangerous command
 - Can easily overwrite valid data (I hope you have backups)
- Because of the power of UPDATE, it is very crucial to use the WHERE clause with it
 - By default, MySQL won't let you update without it
- If you leave off the WHERE, you effectively overwrite an entire column!

UPDATE Syntax

```
UPDATE sometable  
    SET somecol1=value1,  
        somecol2=value2  
    WHERE somefilter;
```

Example:

```
UPDATE patients  
    SET withdrawn = 1  
    WHERE patient_id = 12345;
```

DELETE Your Data

- Are you really sure you should be using this?
 - Just checking
- Even more powerful than `UPDATE`, and often tightly controlled
- The `WHERE` clause is ***REALLY*** important
 - Unless you totally want to dump every row in your table

DELETE Syntax

```
DELETE FROM sometable  
WHERE somefilter;
```

Example:

```
DELETE FROM patients  
WHERE patient_id = 1245;
```

Cascading deletions

- Referential integrity constraints, like foreign keys, often have rules about deletion
- Often, the deletion ***cascades***, or propagates through the database
- For example: you delete a row in your patient table, and all that patient's records are also deleted

Bypassing constraints

- Don't. If they exist, they're probably there for a reason.
- Most RDBMSs will let you bypass them anyway, through various methods
 - The process can be slow - they need to check many constraints, after all
- One which bears mentioning: `TRUNCATE TABLE`


TRUNCATE TABLE

- For those tables that you really hate
- `TRUNCATE TABLE` essentially drops every row of the table instantly
 - Generally, it bypasses checking referential integrity constraints
- Useful if your table does not have any constraints anyway - often quite snappy

CREATE ALTER DROP

- CREATE TABLE makes a new table
- ALTER TABLE changes an existing table
- DROP TABLE, well, drops a table
 - Sounds familiar...

Revisiting the xkcd comic



You should now know a little bit about why this is funny!

Notes on tables

- We will not go into great depth on creating, manipulating, or deleting tables
 - Fairly straightforward in documentation
- Most of this relates to making good data structures, applying constraints, or rules for propagating updates and deletes
- But one thing that is useful...***temporary tables***

Temporary tables

- A *temporary table* is a table which only exists for the duration of your session
 - Disappears when you log out
- You can use these to stage data, store the results of queries, or otherwise manipulate data
 - Can help to make your code shorter!

Creating your temp table

- Temporary tables can be defined using the same basic syntax as regular tables
- Instead of `CREATE TABLE`, you use `CREATE TEMPORARY TABLE`
- This can be combined with `AS` to make easy new temporary tables...

Example Temp Table

```
CREATE TEMPORARY TABLE empno_10004_salaries
  AS
  (
    SELECT *
      FROM salaries
     WHERE emp_no = 10004
  );
```

```
mysql> SELECT * FROM salaries WHERE emp_no = 10004;
```

emp_no	salary	from_date	to_date
10004	40054	1986-12-01	1987-12-01
10004	42283	1987-12-01	1988-11-30
10004	42542	1988-11-30	1989-11-30
10004	46065	1989-11-30	1990-11-30
10004	48271	1990-11-30	1991-11-30
10004	50594	1991-11-30	1992-11-29
10004	52119	1992-11-29	1993-11-29
10004	54693	1993-11-29	1994-11-29
10004	58326	1994-11-29	1995-11-29
10004	60770	1995-11-29	1996-11-28
10004	62566	1996-11-28	1997-11-28
10004	64340	1997-11-28	1998-11-28
10004	67096	1998-11-28	1999-11-28
10004	69722	1999-11-28	2000-11-27
10004	70698	2000-11-27	2001-11-27
10004	74057	2001-11-27	9999-01-01

16 rows in set (0.00 sec)

```
mysql> CREATE TEMPORARY TABLE empno_10004_salaries AS (SELECT * FROM salaries WHERE emp_no = 10004);
```

Query OK, 16 rows affected (0.00 sec)

Records: 16 Duplicates: 0 Warnings: 0

```
mysql> SELECT * FROM empno_10004_salaries;
```

emp_no	salary	from_date	to_date
10004	40054	1986-12-01	1987-12-01
10004	42283	1987-12-01	1988-11-30
10004	42542	1988-11-30	1989-11-30
10004	46065	1989-11-30	1990-11-30
10004	48271	1990-11-30	1991-11-30
10004	50594	1991-11-30	1992-11-29
10004	52119	1992-11-29	1993-11-29
10004	54693	1993-11-29	1994-11-29
10004	58326	1994-11-29	1995-11-29
10004	60770	1995-11-29	1996-11-28
10004	62566	1996-11-28	1997-11-28
10004	64340	1997-11-28	1998-11-28
10004	67096	1998-11-28	1999-11-28
10004	69722	1999-11-28	2000-11-27
10004	70698	2000-11-27	2001-11-27
10004	74057	2001-11-27	9999-01-01

16 rows in set (0.00 sec)

Example Temp Table Results

The power of CREATE + AS

- There are many options when creating tables
 - SQL wants you to define data types, constraints, et cetera...
- `CREATE TABLE foo AS ...` lets you bypass a lot of those steps
 - Useful when you just need a place to stash stuff
- Does not transfer constraints from tables

Getting DDLs

- One of the great things about RDBMSs is that structure is stored with the data
- Data structures are created using something called a ***DDL***, or ***data definition language***
 - In our case, it's just part of SQL
- You can retrieve the DDL statement used to create an object in the database

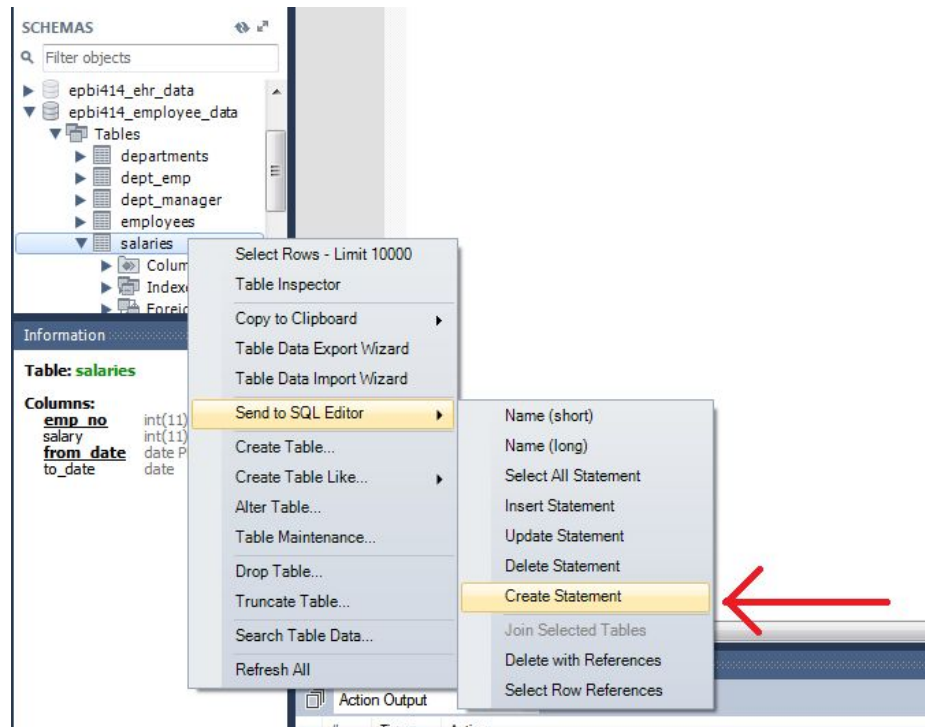
Two methods

- Here are two methods to get the DDL statement used to create an object:
 - `SHOW CREATE TABLE` in command line
(this might be MySQL specific)
 - Using MySQL Workbench

```
mysql> SHOW CREATE TABLE salaries;
+-----+-----+
| Table      | Create Table
+-----+-----+
| salaries | CREATE TABLE `salaries` (
  `emp_no` int(11) NOT NULL,
  `salary` int(11) NOT NULL,
  `from_date` date NOT NULL,
  `to_date` date NOT NULL,
  PRIMARY KEY (`emp_no`,`from_date`),
  CONSTRAINT `salaries_ibfk_1` FOREIGN KEY (`emp_no`) REFERENCES `employees` (`emp_no`) ON DELETE CASCADE
) ENGINE=InnoDB DEFAULT CHARSET=latin1 |
+-----+-----+
1 row in set (0.00 sec)

mysql> |
```

Using the command line



```
Query 1 x
Limit to 10000 rows
1 CREATE TABLE `salaries` (
2   `emp_no` int(11) NOT NULL,
3   `salary` int(11) NOT NULL,
4   `from_date` date NOT NULL,
5   `to_date` date NOT NULL,
6   PRIMARY KEY (`emp_no`,`from_date`),
7   CONSTRAINT `salaries_ibfk_1` FOREIGN KEY (`emp_no`) REFERENCES `employees` (`emp_no`) ON DELETE CASCADE
8 ) ENGINE=InnoDB DEFAULT CHARSET=latin1;
9
```

Using MySQL Workbench

Closing notes, 1

- We are just brushing the tip of what you can do in SQL
 - MySQL is only one RDBMS too...
 - Consider PostgreSQL for a more sophisticated platform
- Your careers will largely start with consuming data in SQL
 - But you never know where you will end up

Closing notes, 2

- SQL can be very fast!
 - Sometimes, it is a lot faster to do work in SQL than it would be in other languages
- SQL is not perfect
 - Just because you can do it in SQL doesn't *necessarily* mean you should
 - Time tradeoff, maintenance tradeoff

Attributions

1. "Is It Worth the Time?" xkcd. <https://xkcd.com/1205/>.
Licensed under Creative Commons Attribution Non-Commerical 2.5. Randall Munroe, author.
2. "Optimization" xkcd. <https://xkcd.com/1691/>.
Licensed under Creative Commons Attribution Non-Commerical 2.5. Randall Munroe, author.
3. "Exploits of a Mom" xkcd. <https://xkcd.com/327/>.
Licensed under Creative Commons Attribution Non-Commerical 2.5. Randall Munroe, author.